

Main Memory

1DV512 – Operating Systems

Dr. Kostiantyn Kucher

`kostiantyn.kucher@lnu.se`

November 18, 2020

Based on the Operating System Concepts slides by Silberschatz, Galvin, and Gagne (2018)

Suggested OSC book complement: Chapter 8

Agenda

- ▶ Motivation and Introduction
- ▶ Contiguous Memory Allocation
- ▶ Paging
- ▶ Swapping
- ▶ Memory Management Examples
- ▶ Summary

Motivation

- ▶ Previously, we discussed multiprogramming and multitasking approaches
⇒ by support multiple processes, we improve the efficiency of computer systems usage
- ▶ However, we must keep many processes in memory ⇒ OS must manage the memory available to the processes!
- ▶ Program must be brought (from disk) into memory and placed within a process for it to be run
- ▶ CPU can only access directly its registers and *main memory* (RAM)
- ▶ *Memory management unit* (MMU) only sees a stream of:
 - ▶ addresses + read requests, and
 - ▶ addresses + data + write requests
- ▶ Register access is done in one CPU clock (or less)
- ▶ Main memory can take many cycles, causing a *stall* ⇒ additional fast *cache* is added between CPU registers and main memory
- ▶ Besides the performance, *protection* of memory is also a concern ⇒ user processes vs the OS, user processes vs each other. . .

Basic Memory Protection

- ▶ OS doesn't usually intervene between the CPU and its memory accesses for performance reasons \Rightarrow some level of hardware support is necessary!
- ▶ Basic protection approach \Rightarrow make sure that each process has a *separate memory space*
- ▶ Typical solution \Rightarrow a pair of *base* and *limit* registers that define the *logical address space* of a process
- ▶ CPU must check every memory access generated in user mode to be sure it is between base and limit for that user
- ▶ The instructions to loading the base and limit registers are privileged \Rightarrow kernel-mode only!




Fig. 9.1 in OSC book




Fig. 9.2 in OSC book

Address Binding

- ▶ Usually, a program resides on a disk as a binary executable file \Rightarrow must be loaded into the memory of a process to be executed!
- ▶ As the process executes, it accesses instructions and data from memory at certain *addresses*
- ▶ *Address binding* of instructions and data to memory addresses can happen at three different stages:
 - ▶ *Compile time* \Rightarrow if memory location known a priori, *absolute code* can be generated
 - ▶ *Load time* \Rightarrow must generate *relocatable code* if memory location is not known at compile time
 - ▶ *Execution time* \Rightarrow binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ▶ Need hardware support for address maps (e.g., base and limit registers)
 - ▶ Most operating systems use this method!




Fig. 9.3 in OSC book

Logical and Physical Addresses

- ▶ *Logical address* \Rightarrow generated by the CPU; also referred to as *virtual address*
- ▶ *Physical address* \Rightarrow address seen by the memory management unit (MMU)
- ▶ Logical and physical addresses are the same in compile-time and load-time address-binding schemes, and different in the execution-time address binding
- ▶ One simple approach for execution-time *relocation* for MMU \Rightarrow add the value from the *relocation register* to each logical address
- ▶ The user program deals with logical addresses (in range $0 \dots \text{max}$) \Rightarrow it never sees the real physical addresses (in range $R + 0 \dots R + \text{max}$)!
- ▶ Thus, a logical address space is bound to a separate physical address




Fig. 9.4 in OSC book




Fig. 9.5 in OSC book

Dynamic Loading and Dynamic Linking

- ▶ *Dynamic loading* \Rightarrow instead of loading the complete program in memory, only the main necessary part (routine) is loaded
- ▶ All routines kept on disk in relocatable load format and loaded as they are invoked for the first time
- ▶ No special support from the OS is required \Rightarrow however, OS can assist by providing libraries/tools to implement this
- ▶ *Dynamic linking* \Rightarrow binding delayed until the execution time, so that some libraries could be *shared* (another term: *dynamically loaded libraries*, DLLs)
- ▶ When a corresponding routine is invoked, the OS loads the library into memory (if necessary) and adjusts the addresses in the calling program
- ▶ Dynamic linking and shared libraries typically require support from the OS

Agenda

- ▶ Motivation and Introduction
- ▶ **Contiguous Memory Allocation**
- ▶ Paging
- ▶ Swapping
- ▶ Memory Management Examples
- ▶ Summary

Contiguous Allocation

- ▶ Main memory must support both OS and user processes
- ▶ Limited resource \Rightarrow must be allocate efficiently!
- ▶ *Contiguous allocation* is one early and straightfotward method
- ▶ Main memory usually divided into two partitions:
 - ▶ Resident operating system \Rightarrow usually held in high memory with interrupt vector
 - ▶ User processes then held in low memory \Rightarrow each process contained in single *contiguous* section of memory
 - ▶ Note: placing OS in high memory is a design choice, and low memory could be used instead in some cases

Memory Protection

- ▶ Relocation register + limit register can be used to protect user processes from each other, and from changing operating-system code and data
- ▶ Base register contains the value of the smallest physical address
- ▶ Limit register contains the range of logical addresses \Rightarrow each logical address must be less than the limit register
- ▶ MMU maps logical address dynamically \Rightarrow can allow actions such as kernel code being transient and kernel changing size




Fig. 9.6 in OSC book

Memory Allocation and Variable Partitions

- ▶ *Variable-partition allocation* \Rightarrow assign a memory partition to each process according to its needs
- ▶ *Hole* \Rightarrow a block of available memory; holes of various size are scattered throughout memory
- ▶ When a process arrives, it is allocated memory from a hole large enough to accommodate it
- ▶ An exiting process frees its partition, and adjacent free partitions are combined
- ▶ Operating system maintains information about: a) allocated partitions and b) free partitions (holes)




Fig. 9.7 in OSC book

Dynamic Storage Allocation Problem and Fragmentation


- ▶ How to satisfy a request of size n from the list of free holes?
 - ▶ *First fit* \Rightarrow allocate the first hole that is big enough
 - ▶ *Best fit* \Rightarrow allocate the smallest hole that is big enough; must search entire list, unless it is ordered by size
 - ▶ *Worst fit* \Rightarrow allocate the largest hole; must also search entire list, unless it is ordered by size
- ▶ According to simulations, *first fit* and *best fit* are both good options with regard to storage utilization; and *first fit* is generally faster
- ▶ Both strategies suffer from *external fragmentation* \Rightarrow enough total memory space exists to satisfy the next request, but it is not contiguous!
 - ▶ *First fit* analysis reveals that for N allocated blocks, $0.5N$ is lost to fragmentation \Rightarrow *50-percent rule*
 - ▶ Can be addressed by *compaction* (not always possible) or *paging* \Rightarrow discussed later!
- ▶ *Internal fragmentation* \Rightarrow allocated memory (multiple fixed-size blocks) may be slightly larger than requested memory

Agenda

- ▶ Motivation and Introduction
- ▶ Contiguous Memory Allocation
- ▶ **Paging**
- ▶ Swapping
- ▶ Memory Management Examples
- ▶ Summary

Paging — Basic Method

- ▶ *Paging* \Rightarrow a memory management scheme that permits a process's physical address space to be non-contiguous
- ▶ Used in most modern OSs and implemented with both software (OS) and hardware support
- ▶ Prevents external fragmentation, although internal fragmentation is still possible
- ▶ Basic method:
 - ▶ Divide physical memory into fixed-sized blocks called *frames*
 - ▶ Divide logical memory into blocks of same size called *pages*
 - ▶ Frame size \Rightarrow usually power of 2, between 512 bytes and 16 MB
 - ▶ Logical address space can now be much larger than the amount of RAM!
 - ▶ For a program with a size of N pages $\Rightarrow N$ free frames required in total
 - ▶ Every address generated by the CPU is divided into two parts: a *page number* (p) and a *page offset* (d) for given logical address space of size 2^m and page size 2^n



Sec. 9.3.1 in OSC book

Page Table

- ▶ A *page table* is used to translate logical to physical addresses for a process
- ▶ It contains the base address of each frame in physical memory
- ▶ Thus, a page number p and a page offset d are used with the page table
- ▶ MMU replaces p in the logical address with the corresponding frame number f ; d is not changed




Fig. 9.8 in OSC book

Paging Model of Logical and Physical Memory

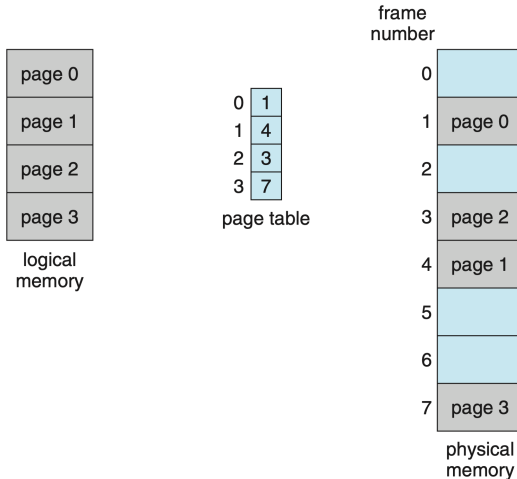


Fig. 9.9 in OSC book

Paging Example

- Here, $n = 2$ bits \Rightarrow page size is $2^n = 4$ bytes
- $m = 4$ bits \Rightarrow logical address space size is $2^m = 16$ bytes
- Available physical memory is 32 bytes \Rightarrow 8 pages
- For logical address 0 $\Rightarrow p = 0, d = 0 \Rightarrow f = 5 \Rightarrow$ physical address is $f \cdot 2^n + d = 5 \cdot 4 + 0 = 20$
- For logical address 13 $\Rightarrow p = 3, d = 1 \Rightarrow f = 2 \Rightarrow$ physical address is $f \cdot 2^n + d = 2 \cdot 4 + 1 = 9$

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

Fig. 9.10 in OSC book

Free Frames Before and After Process Allocation




Fig. 9.11 in OSC book

Hardware Support

- ▶ Page tables are *per-process* data structures \Rightarrow stored in PCB for each process
- ▶ Hardware support approaches for improving page table processing performance:
 - ▶ Dedicated high-speed CPU registers \Rightarrow efficient address translation, but increased context switching time + not feasible for thousands of pages!
 - ▶ *Page-table base register (PTBR)* \Rightarrow keep a pointer to respective the page table stored in main memory
 - ▶ *Translation look-aside buffer (TLB)* \Rightarrow small and high-speed hardware cache for recent look-ups




Fig. 9.12 in OSC book

Memory Protection

- ▶ Memory protection implemented by associating protection bit with each frame \Rightarrow indicate if *read-only* or *read-write* access is allowed
- ▶ Can also add more bits to indicate page *execute-only*, and so on
- ▶ *Valid-invalid* bit attached to each entry in the page table:
 - ▶ *valid* \Rightarrow the page is in the process logical address space, and is thus a legal page
 - ▶ *invalid* \Rightarrow the page is not in the process logical address space
 - ▶ Alternatively, use *page-table length register (PTLR)*
- ▶ Any violations result in a trap to the kernel

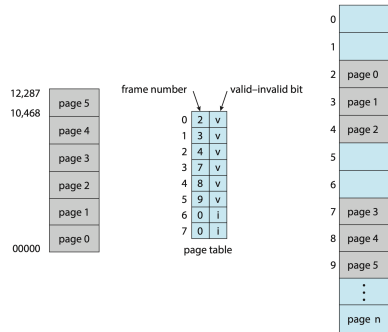


Fig. 9.13 in OSC book

Shared Pages

- ▶ *Shared code* \Rightarrow only one copy of read-only (*reentrant*) code shared among processes
- ▶ Efficient memory space usage \Rightarrow e.g., for widely used shared libraries such as standard C library
- ▶ Also useful for IPC if sharing of read-write pages is allowed
- ▶ Each process also keeps its own private code and data
- ▶ The pages for the private code and data can appear anywhere in the logical address space




Fig. 9.14 in OSC book

Advanced Page Table Structures

- ▶ Memory structures for paging can get huge using basic, straightforward methods \Rightarrow millions of page table entries, megabytes of memory for each process for the page table itself. . .
- ▶ Thus, several more advanced approaches are applied \Rightarrow the idea is to divide the page table into smaller units
- ▶ *Hierarchical Paging* \Rightarrow multiple page tables (e.g., two-level) + paging the page table (so-called *forward-mapped page table*)
- ▶ *Hashed Page Tables* \Rightarrow for >32-bit systems, more efficient to use a virtual page number hashed into a page table
- ▶ *Inverted Page Tables* \Rightarrow keep only a single page table and track all physical pages/frames used by processes




Fig. 9.15 in OSC book




Fig. 9.17 in OSC book




Fig. 9.18 in OSC book

Agenda

- ▶ Motivation and Introduction
- ▶ Contiguous Memory Allocation
- ▶ Paging
- ▶ **Swapping**
- ▶ Memory Management Examples
- ▶ Summary

Swapping

- ▶ If the amount of physical memory is insufficient \Rightarrow a process can be *swapped* temporarily out of memory to a *backing store*, and then brought back into memory for continued execution
- ▶ *Backing store* \Rightarrow fast disk large enough to accommodate copies of all memory images for all users
- ▶ Standard swapping was used in traditional Unix systems, but is no longer used in modern OSs \Rightarrow too slow to swap complete processes!
- ▶ E.g., mobile OSs (iOS, Android) remove process data if free memory is low
- ▶ Most modern systems support *swapping with paging* \Rightarrow the *page out* operation moves a page to the backing store, and the *page in* operation does the opposite




Fig. 9.19 in OSC book




Fig. 9.20 in OSC book

Agenda

- ▶ Motivation and Introduction
- ▶ Contiguous Memory Allocation
- ▶ Paging
- ▶ Swapping
- ▶ Memory Management Examples
- ▶ Summary

Intel 32- and 64-bit Architectures

- ▶ IA-32 involves *segmentation* + paging
 \Rightarrow up to 16000 segments per process, each segment up to 4 GB \Rightarrow based on the logical address, a 32-bit *linear address* is generated
- ▶ Logical address space for each process divided into *local* (process-private) and *shared* partitions
- ▶ Two-level paging involving either 4 KB or 4 MB pages
- ▶ *Page Address Extension (PAE)* used to address >4 GB \Rightarrow using a three-level scheme leads to the support of up to 64 GB of RAM, if supported by OS
- ▶ *x86-64* supports up to 64-bit address space \Rightarrow 48 bits are currently used in practice, with a four-level paging scheme



Fig. 9.21 in OSC book




Fig. 9.23 in OSC book




Fig. 9.24 in OSC book




Fig. 9.25 in OSC book

ARM 64-bit Architecture

- ▶ *ARMv8* supports three different *translation granules*: 4 KB, 16 KB, and 64 KB
- ▶ Each translation granule provides different page sizes, as well as larger sections of contiguous memory, known as *regions*
- ▶ 4 KB and 16 KB granules \Rightarrow up to four levels of paging; 64 KB \Rightarrow up to three
- ▶ 48 (out of 64) bits currently used for addresses
- ▶ *Translation Table Base Register (TTBR)* \Rightarrow points to table level 0 for the current thread
- ▶ ARM supports two levels of translation look-aside buffers (TLBs) \Rightarrow two *micro TLBs* (instructions + data) at inner level and one *main TLB* at outer level

Translation Granule Size	Page Size	Region Size
4 KB	4 KB	2 MB, 1 GB
16 KB	16 KB	32 MB
64 KB	64 KB	512 MB

Sec. 9.7 in OSC book



Fig. 9.26 in OSC book




Fig. 9.27 in OSC book

Agenda

- ▶ Motivation and Introduction
- ▶ Contiguous Memory Allocation
- ▶ Paging
- ▶ Swapping
- ▶ Memory Management Examples
- ▶ **Summary**

Summary

- ▶ One way to allocate memory address space to each process is through the use of base and limit registers
- ▶ CPU generates a logical address, and MMU translates it to a physical address
- ▶ Modern OSs typically use memory paging with the page size of either 4 KB or 8 KB, and additionally support *huge pages* for special scenarios
- ▶ 64-bit and even 32-bit architectures often require advanced paging techniques involving multiple paging levels and hashing