

# Processes and Threads

*1DV512 – Operating Systems*

Dr. Kostiantyn Kucher

`kostiantyn.kucher@lnu.se`

November 10, 2020

Based on the Operating System Concepts slides by Silberschatz, Galvin, and Gagne (2018)

Suggested OSC book complement: Chapters 3 & 4

# Agenda

- ▶ Motivation and Introduction
- ▶ Processes
- ▶ Process Operations
- ▶ Threads
- ▶ Summary

# Motivation

- ▶ Typically, we expect *general-purpose* computer systems to be designed for more than a single specific program. . .
  - ▶ E.g., imagine a computer that can *only* run a “Hello world” program
  - ▶ Notice the focus on *general-purpose*  $\Rightarrow$  specialized, albeit more restricted designs can be motivated for microcontrollers, etc.
- ▶ Early computers  $\Rightarrow$  a single program  $\Rightarrow$  a single user. . .
- ▶ Later, the *multiprogramming* approach  $\Rightarrow$  batches of *jobs* (code + data)  $\Rightarrow$  multiple users (but they have to wait. . . )
- ▶ Later, the *multitasking* approach  $\Rightarrow$  frequently switching between multiple *tasks*  $\Rightarrow$  multiple users able to work interactively!
- ▶ *Process* is a more recent and more general term describing a program in *execution*




Fig. 1.12 in OSC book

# Agenda

- ▶ Motivation and Introduction
- ▶ **Processes**
- ▶ Process Operations
- ▶ Threads
- ▶ Summary

# Structure of a Process

- ▶ A *program* is a passive entity stored on disk (an *executable file*)  $\Rightarrow$  a *process* is active
- ▶ A program becomes a process when an executable file is loaded into memory (e.g., by clicking an app icon. . .)
- ▶ Same program can be launched several times (or by several users)  $\Rightarrow$  multiple processes!
- ▶ Process  $\Rightarrow$  memory contents + CPU register contents + CPU program counter state
- ▶ Typical structure of a process in memory:
  - ▶ *Text section*  $\Rightarrow$  executable code (fixed size!)
  - ▶ *Data section*  $\Rightarrow$  global variables (fixed size!)
  - ▶ *Heap*  $\Rightarrow$  dynamically allocated memory
  - ▶ *Stack*  $\Rightarrow$  temporary data / *activation records* for functions (local variables, arguments, return addresses. . .)




Fig. 3.1 in OSC book

# Process States




Fig. 3.2 in OSC book

As a process executes, it changes *state*, typically in the following way:

- ▶ *New* ⇒ The process is being created
- ▶ *Ready* ⇒ The process is waiting to be assigned to a processor
- ▶ *Running* ⇒ Instructions are being executed
- ▶ *Waiting* ⇒ The process is waiting for some event to occur
- ▶ *Terminated* ⇒ The process has finished execution

# Process Control Block

Each process is represented with a *process/task control block*:

- ▶ Process state  $\Rightarrow$  ready, running, etc.
- ▶ Process number / PID
- ▶ CPU registers  $\Rightarrow$  contents of all process-centric registers
- ▶ CPU scheduling information  $\Rightarrow$  priorities, scheduling queue pointers
- ▶ Memory-management information  $\Rightarrow$  memory allocated to the process
- ▶ Accounting information  $\Rightarrow$  CPU used, clock time elapsed since start, time limits, etc.
- ▶ I/O status information  $\Rightarrow$  I/O devices allocated to process, list of open files
- ▶ *Multiple program counters per process*  $\Rightarrow$  possibility of multiple control *threads*!




Fig. 3.3 in OSC book

# Process Scheduling




Fig. 3.4 in OSC book

- ▶ *Process scheduler* selects available processes for execution on a CPU core and maintains *scheduling queues*
- ▶ *Ready queue*  $\Rightarrow$  set of all processes residing in main memory, ready, and waiting to execute
- ▶ *Wait queues*  $\Rightarrow$  sets of processes waiting for particular events
- ▶ A process goes back and forth through the queues and CPU execution until it terminates
- ▶ Further details on scheduling  $\Rightarrow$  next lecture!



# Queueing Diagram Representation

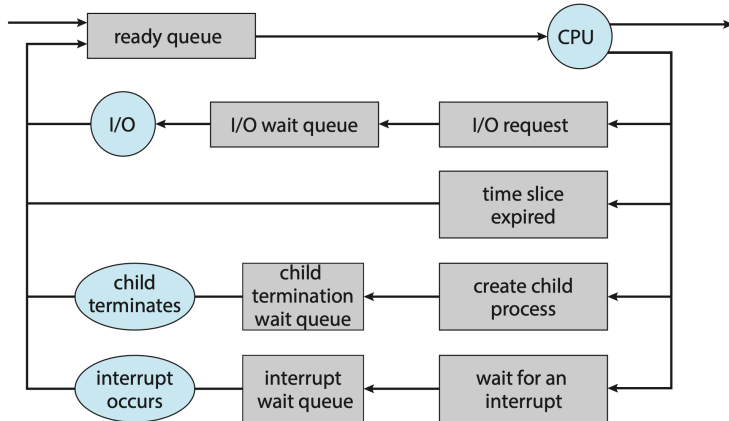


Fig. 3.5 in OSC book

# Context Switching

When switching between several processes, the *context* information must be saved  
 $\Rightarrow$  CPU register values + process state + memory management information

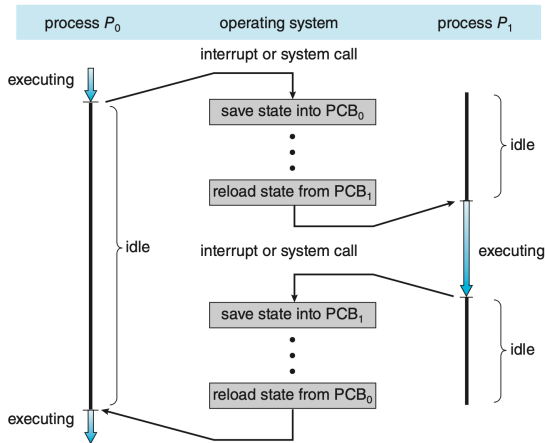


Fig. 3.6 in OSC book

# Agenda

- ▶ Motivation and Introduction
- ▶ Processes
- ▶ **Process Operations**
- ▶ Threads
- ▶ Summary

# Typical Operations and The Process Tree

- ▶ Processes are typically identified in a unique way  $\Rightarrow$  *process identifier (pid)*
- ▶ OS typically support at least *creation* and *termination* of processes
- ▶ A process is typically created by another, parent process  $\Rightarrow$  the *process tree*
- ▶ In Unix-like systems:
  - ▶ pid 0  $\Rightarrow$  scheduler (part of kernel rather than user process, not always included in the process tree)
  - ▶ pid 1  $\Rightarrow$  the *init* process (in recent Linux distributions, *systemd* instead)  $\Rightarrow$  parent for other processes
  - ▶ Other parts of modern kernels can be exposed with pids, too, e.g., *[kthreadd]*, *[ksoftirqd]*, etc.




Fig. 3.7 in OSC book

# Design Choices for Process Creation

- ▶ *Resource sharing* options:
  - ▶ Parent and children share all resources
  - ▶ Children share a subset of parent's resources
  - ▶ Parent and children share no resources
  
- ▶ *Execution* options:
  - ▶ Parent and children execute concurrently
  - ▶ Parent waits until children terminate
  
- ▶ *Address space* options:
  - ▶ Child is created as a duplicate of parent
  - ▶ Child is created with a specific program loaded

# Process Creation in Unix-like Systems

- ▶ `fork()` system call creates new process  $\Rightarrow$  copies the memory space of the parent, and both parent and child continue to execute!
- ▶ To differentiate between such “clone” processes  $\Rightarrow$  use the return value of `fork()`:
  - ▶  $<0 \Rightarrow$  error
  - ▶  $=0 \Rightarrow$  executing within the new, child process
  - ▶  $>0 \Rightarrow$  executing within the old, parent process  $\Rightarrow$  pid of the child process!
- ▶ To load and start executing a different program from disk  $\Rightarrow$  invoke the `exec()` system call afterwards  $\Rightarrow$  memory space replaced completely
- ▶ In Windows  $\Rightarrow$  a different program has to be specified immediately for `CreateProcess()`




Fig. 3.9 in OSC book

# Process Termination in Unix-like Systems

- ▶ `wait()` system call  $\Rightarrow$  wait (suspend execution) until *one* of the children terminates
- ▶ `waitpid()` system call  $\Rightarrow$  wait (suspend execution) until any or some (= specific pid or *process group*) child terminates
- ▶ `exit()` system call (from child)  $\Rightarrow$  terminate, return a status value to the parent, and release resources (memory, files, etc.)
- ▶ `abort()` system call (from parent)  $\Rightarrow$  request termination of a child process
- ▶ Parent process did not invoke `wait()` yet  $\Rightarrow$  the child process is a *zombie*
- ▶ Parent process terminated without invoking `wait()`  $\Rightarrow$  the child process is an *orphan*

# Agenda

- ▶ Motivation and Introduction
- ▶ Processes
- ▶ Process Operations
- ▶ **Threads**
- ▶ Summary



# Single- and Multithreaded Processes

- ▶ *Process*  $\Rightarrow$  process ID, register set, stack, program counter, code section, data section, open files, ...
- ▶ *Thread*  $\Rightarrow$  more lightweight abstraction of a computational task  $\Rightarrow$  thread ID, register set, stack, program counter.  $\Rightarrow$  code and most data shared within the same process
- ▶ Traditional approach  $\Rightarrow$  a single thread of execution per process
- ▶ Later approaches  $\Rightarrow$  support for multithreaded processes

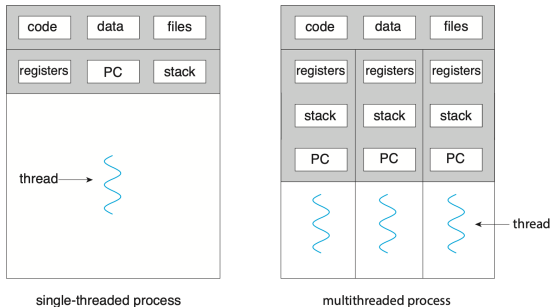


Fig. 4.1 in OSC book

# Concurrency and Parallelism




Fig. 4.3 in OSC book




Fig. 4.4 in OSC book

# Data and Task Parallelism

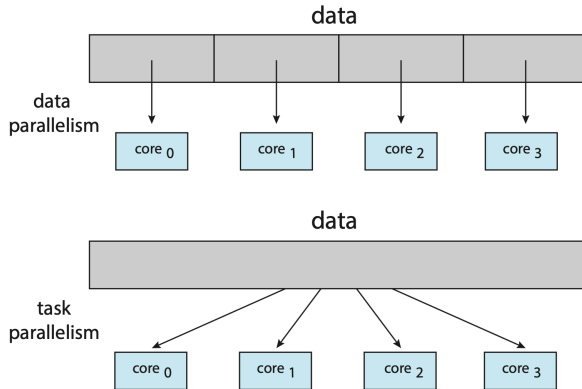


Fig. 4.5 in OSC book

# Multithreading Models

There are multiple models depending on whether the OS supports *kernel threads*, and how *user threads* are mapped to them:

- ▶ *Many-to-one* model  $\Rightarrow$  simple thread management, but failing to make use of multicore hardware
- ▶ *One-to-one* model  $\Rightarrow$  reasonable, but with overheads for creating kernel threads
- ▶ *Many-to-many* model  $\Rightarrow$  flexible, but difficult to implement
- ▶ *Two-level* model  $\Rightarrow$  adapted variation of the many-to-many model




Fig. 4.6 in OSC book




Fig. 4.7 in OSC book




Fig. 4.8 in OSC book




Fig. 4.9 in OSC book




Fig. 4.10 in OSC book

# Agenda

- ▶ Motivation and Introduction
- ▶ Processes
- ▶ Process Operations
- ▶ Threads
- ▶ Summary

# Summary

- ▶ Process  $\Rightarrow$  program in execution, including the data, resources, code, and program counter
- ▶ Based on the OS, processes change their states during their life cycles, and are typically related to each other with parent-child relationships
- ▶ Process creation paradigm in Unix-like systems  $\Rightarrow$  `fork()` + `exec()`
- ▶ Thread  $\Rightarrow$  more lightweight abstraction than a process, with several threads sharing the same code, data, and resources
- ▶ Multithreading can improve concurrency (and thus computational efficiency), especially for multicore hardware
- ▶ More on scheduling and process & thread synchronization issues  $\Rightarrow$  following lectures!