

Synchronization and Deadlocks

1DV512 – Operating Systems

Dr. Kostiantyn Kucher

`kostiantyn.kucher@lnu.se`

November 17, 2020

Based on the Operating System Concepts slides by Silberschatz, Galvin, and Gagne (2018)

Suggested OSC book complement: Chapters (3 & 6 & 7

Agenda

- ▶ Motivation and Introduction
- ▶ Inter-Process Communication
- ▶ Process and Thread Synchronization Approaches
- ▶ Deadlocks
- ▶ Summary

Motivation

- ▶ We have previously established the concepts of a *process* and a *thread*
- ▶ We have also discussed how such tasks can be *scheduled* for execution to proceed with their computations
- ▶ What if the computations depend on the results or intermediate data from another process or thread?
 - ▶ We discussed how a parent process can check the exit status of the child process, but there is surely more to this topic!
 - ▶ A *lot* of processes and threads can be executed concurrently by the OS
⇒ concurrent access to shared data may result in data inconsistency
- ▶ We must discuss the models for exchanging data between the tasks, and preventing problematic situations that might arise

Agenda

- ▶ Motivation and Introduction
- ▶ **Inter-Process Communication**
- ▶ Process and Thread Synchronization Approaches
- ▶ Deadlocks
- ▶ Summary

Main Models of Inter-Process Communication

- ▶ Processes within a system may be *independent* or *cooperating*
- ▶ Cooperating process can affect or be affected by other processes, including sharing data
- ▶ Reasons for cooperating processes:
 - ▶ Information sharing
 - ▶ Computation speedup
 - ▶ Modularity and convenience
- ▶ Cooperating processes need *interprocess communication (IPC)*
- ▶ Two main models of IPC:
 - ▶ *Shared memory* \Rightarrow faster for common access to larger amounts of data
 - ▶ *Message passing* \Rightarrow easier and more efficient for smaller amounts of data
- ▶ Major issue: how to provide a mechanism that will allow the user processes to *synchronize*




Fig. 3.11 in OSC book

Producer-Consumer Problem

- ▶ Common paradigm for cooperating processes \Rightarrow *producer* process produces information that is consumed by a *consumer* process
- ▶ The producer and consumer must be synchronized, so that the consumer does not try to consume an item that has not yet been produced
- ▶ One possible solution for the producer-consumer problem \Rightarrow shared memory
- ▶ The communication is under the control of the user processes, not the operating system
- ▶ A buffer of items that can be filled by the producer and emptied by the consumer, with two main variations:
 - ▶ *unbounded-buffer* places no practical limit on the size of the buffer
 - ▶ Producer never waits
 - ▶ Consumer waits if there is no buffer to consume
 - ▶ *bounded-buffer* assumes that there is a fixed buffer size
 - ▶ Producer must wait if all buffers are full
 - ▶ Consumer waits if there is no buffer to consume

Producer-Consumer Problem (cont.)

- ▶ The shared buffer is implemented as a circular array with two logical pointers: `in` and `out`
 - ▶ `in` points to the next free position in the buffer
 - ▶ `out` points to the first full position in the buffer
- ▶ The buffer is full when $((in + 1) \% BUFFER_SIZE) == out$, and empty when `in == out`
- ▶ Simpler solution allows the usage of `BUFFER.SIZE-1` elements
- ▶ Solution using all `BUFFER.SIZE` elements can lead to *race conditions* if both processes try to access the buffer simultaneously \Rightarrow more on synchronization later!

```
item buffer[BUFFER.SIZE];
int in = 0;
int out = 0;
```

Sec. 3.5 in OSC book

```
item next_produced;

while (true) {
    /* produce an item in next_produced */

    while (((in + 1) \% BUFFER.SIZE) == out)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) \% BUFFER.SIZE;
}
```

Fig. 3.12 in OSC book

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) \% BUFFER.SIZE;

    /* consume the item in next_consumed */
}
```

Fig. 3.13 in OSC book

Message Passing

- ▶ Message passing is an alternative to shared memory \Rightarrow support for `send(message)` and `receive(message)` operations, with either fixed or variable message size
- ▶ *Direct* communication \Rightarrow a pair of processes name each other explicitly and send messages to each other (can also be unidirectional)
- ▶ *Indirect* communication \Rightarrow messages are directed and received from mailboxes (also referred to as *ports*)
- ▶ Message passing may be either *blocking* (synchronous) or *non-blocking* (asynchronous)
- ▶ Messages reside in a temporary buffering queue of *zero* capacity (\Rightarrow sender must wait for receiver), *bounded*/finite capacity (\Rightarrow sender must wait if full queue), or *unbounded* capacity (\Rightarrow sender never waits)
- ▶ The producer-consumer problem can be approached with message passing, too:

```
message next_produced;

while (true) {
    /* produce an item in next_produced */
    send(next_produced);
}
```

Fig. 3.15 in OSC book

```
message next_consumed;

while (true) {
    receive(next_consumed);
    /* consume the item in next_consumed */
}
```

Fig. 3.15 in OSC book

IPC Examples

- ▶ Shared memory
- ▶ Pipes
 - ▶ Once established, reading and writing operations similar to regular file read/write API
 - ▶ Can be typically uni- or bidirectional
 - ▶ Unidirectional (*ordinary* or *anonymous*) pipes ⇒ a pair of pipes might be necessary to communicate between a parent and a child process
 - ▶ Biredictional (*named*) pipes ⇒ often visible as special files in the file system, can be used by multiple processes
- ▶ Sockets
 - ▶ Communication between multiple *hosts* (⇒ *network sockets*) or as an efficient IPC tool within the same host (⇒ *Unix domain sockets*)
- ▶ Remote Procedure Calls (RPC)




Fig. 3.20 in OSC book




Fig. 3.26 in OSC book




Fig. 3.29 in OSC book

Agenda

- ▶ Motivation and Introduction
- ▶ Inter-Process Communication
- ▶ Process and Thread Synchronization Approaches
- ▶ Deadlocks
- ▶ Summary

Synchronization and The Critical Section Problem

- ▶ Previously, we saw how concurrent access and manipulation of shared data in the consumer-producer problem can lead to inconsistency issues \Rightarrow *race conditions*
- ▶ One scenario: several processes invoking `fork()` at the *same* time and receiving the same child PID!
- ▶ *Synchronization* issues even more important, given the prominence of modern multicore systems with support for multithreading




Fig. 6.2 in OSC book

Critical Section Problem

- ▶ Critical section \Rightarrow a section of code where the data shared with another process is accessed or updated
- ▶ While one process is in the critical section \Rightarrow no other processes allowed to enter their critical section!
- ▶ Each process must ask permission to enter critical section in *entry section*, may follow critical section with *exit section*, then *remainder section*
- ▶ Requirements for possible solution designs:
 - ▶ *Mutual exclusion*
 - ▶ *Progress* \Rightarrow if several processes are waiting to enter the critical section, one of them should eventually be able to do it!
 - ▶ *Bounded waiting*
- ▶ The problem is very important for kernel-mode code
- ▶ Solutions can be completely software-based (\Rightarrow often not sufficient for multicore systems) or might involve support from hardware (\Rightarrow *memory barriers/fences*, *atomic instructions* such as *test and set* or *compare and swap*, and *atomic variables*)




Fig. 6.1 in OSC book

```

boolean test_and_set(boolean *target) {
    boolean rv = *target;
    *target = true;

    return rv;
}

```

Fig. 6.5 in OSC book

```

int compare_and_swap(int *value, int expected, int new_value) {
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}

```

Fig. 6.7 in OSC book

Mutex Lock

- ▶ OS designers typically build software tools to solve critical section problem
- ▶ Simplest is *mutex lock* (using the mutual exclusion concept) \Rightarrow boolean variable indicating if lock is available or not
- ▶ Provides `acquire()` and `release()` operations \Rightarrow must be atomic! (usually implemented using atomic hardware instructions)
- ▶ If a process is forced to loop (*spin*) a lot while waiting \Rightarrow such mutex is also called a *spinlock*
 - ▶ Can be preferable, if the waiting time is shorter than the combined time of context switches!

```

while (true) {
    acquire lock
    critical section
    release lock
    remainder section
}

```

Fig. 6.10 in OSC book

```

do {
    while (test_and_set(&lock))
        ; /* do nothing */

    /* critical section */

    lock = false;

    /* remainder section */
} while (true);

```

Fig. 6.6 in OSC book

```

while (true) {
    while (compare_and_swap(&lock, 0, 1) != 0)
        ; /* do nothing */

    /* critical section */

    lock = 0;

    /* remainder section */
}

```

Fig. 6.8 in OSC book

Semaphore

- ▶ *Semaphore* \Rightarrow synchronization tool that provides more sophisticated ways (than mutex locks) for processes to synchronize their activities
- ▶ Relies on an integer counter value S that can only be accessed using two atomic operations, `wait()` and `signal()`
- ▶ *Binary* semaphore $\Rightarrow S$ is either 0 or 1
 \Rightarrow equivalent to a mutex lock
- ▶ *Counting* semaphore $\Rightarrow S$ has an unrestricted range
- ▶ Powerful tool, but no guarantees that the users' code will use it correctly!

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

Sec. 6.6 in OSC book

```
signal(S) {  
    S++;  
}
```

Sec. 6.6 in OSC book

```
S1;  
signal(synch);
```

Sec. 6.6.1 in OSC book

```
wait(synch);  
S2;
```

Sec. 6.6.1 in OSC book

Monitor

- ▶ *Monitor* \Rightarrow an high-level abstraction that provides a convenient and effective mechanism for process synchronization
- ▶ Abstract data type, internal variables only accessible by code within the procedure
- ▶ Only one process may be active within the monitor at a time




Fig. 6.12 in OSC book

Agenda

- ▶ Motivation and Introduction
- ▶ Inter-Process Communication
- ▶ Process and Thread Synchronization Approaches
- ▶ **Deadlocks**
- ▶ Summary

Deadlocks

- ▶ A system consists of a finite number of resources to be distributed among a number of competing processes or threads
- ▶ Each process normally utilizes a resource as follows:
 - ▶ request
 - ▶ use
 - ▶ release
- ▶ *Deadlock* \Rightarrow every task is waiting for an event that can be caused only by another task

```
/* thread.one runs in this function */
void *do_work.one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}
```

```
/* thread.two runs in this function */
void *do_work.two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

Fig. 8.01 in OSC book

Deadlock Conditions

- ▶ Deadlock can arise if four conditions hold simultaneously:
 - ▶ Mutual exclusion: only one process at a time can use a resource
 - ▶ Hold and wait: a process holding at least one resource is waiting to acquire additional resources held by other processes
 - ▶ No preemption: a resource can be released only voluntarily by the process holding it, after that process has completed its task
 - ▶ Circular wait: there is a cycle of processes waiting for each other's actions

Deadlock Handling Approaches

- ▶ Ignore the problem and pretend that deadlocks never occur in the system
- ▶ Ensure that the system will never enter a deadlock state:
 - ▶ Deadlock prevention
 - ▶ Deadlock avoidance
- ▶ Allow the system to enter a deadlock state and then recover

Agenda

- ▶ Motivation and Introduction
- ▶ Inter-Process Communication
- ▶ Process and Thread Synchronization Approaches
- ▶ Deadlocks
- ▶ **Summary**

Summary

- ▶ The typical models for inter-process communication involve shared memory and message passing \Rightarrow synchronization issues often exist
- ▶ Support for synchronization between processes and threads may involve software and hardware aspects
- ▶ The most typical approaches for synchronization include mutex locks and semaphores
- ▶ Deadlock situations can emerge based on the patterns of access to the resources
- ▶ **Additional demos** (source code + slides) on thread programming with Java will be published on MyMoodle \Rightarrow incl. thread synchronization approaches